

CSC 488S/CSC 2107S Lecture Notes

These lecture notes are provided for the personal use of students taking CSC488H1S or CSC2107S in the

Winter 2012/2013 term at the University of Toronto

Copying for purposes other than this use and all forms of distribution are expressly prohibited.

©David B. Wortman, 2002,2003,2004,2008,2009,2010,2012,2013

©Marsha Chechik, 2005,2006,2007

0

Dynamic Storage Design Issues

- Roots and pointer finding.
Is it possible to find all pointers (roots) to allocated objects?
This is really hard to get correct.
- Mobility of data. It is only possible to move a data object if it is possible to (eventually) update *all pointers* to the object to point at the new location.
- Immediacy of storage reclamation. Is recyclable memory made available immediately or only periodically.
- Program tolerance to interruption and/or arbitrary delay.
Can mutation and collection proceed concurrently?
- Processing Cost – Time and Space Overhead
- Fragmentation of heap memory.
Performance Degradation as heap utilization increases.
- Recovering cyclic data structures.

313

Dynamic Memory Management

- Management of dynamically created runtime storage is a central issue in the implementation of many programming languages.
- Most implementations allocate storage in a *heap* with some mechanism for reclaiming memory that is no longer in use.
- **Garbage collection**^a is a tool for automatically managing dynamically allocated memory.
 - Identify dynamically allocated memory that is no longer in use.
 - Recycle that memory for later reuse
- Some programming language require Garbage Collection (Java, Lisp), many others allow it.
- The presence of Garbage Collection adds runtime overhead and may introduce unpredictable *stalls* to program execution

^aRichard Jones, Antony Hoskin, Elliot Moss, *The Garbage Collection Handbook: The Art of Automatic Memory Management*, Chapman & Hall, 2011

312

Basic Principle

- The key concept for garbage collection is **reachability**. A program can only access some object if it has a memory address (i.e. pointer) for the object. An object that has been allocated but is no longer reachable, is *garbage* and may be recycled.
- *Roots* are base locations that may hold the address of dynamically allocated objects. Possible roots include: processor registers, all kinds of variables and temporaries in activation records.
- Dynamically allocated memory can only be accessed through some chain of pointers starting with some root. An object is *reachable* if and only if there is some path from a root to the object.
- Two separate issues:
 - identifying reachable objects
 - managing and recycling memory efficiently.

314

Classic Dynamic Memory Design Approaches

- *Reference Counting* - Every allocated object contains a *reference count*. This count is incremented when a new pointer to the object is created and decremented when a pointer to the object is destroyed. A reference count of zero implies the object has no pointers pointing at it and so it can be reclaimed. Incremental algorithm.
- *Mark and Sweep Collection* - A garbage collector algorithm starts at all roots recursively marking all nodes that can be reached. When this algorithm finishes, any nodes that are not marked can be reclaimed.

Conservative marking, takes more space, but use a simpler algorithm to identify pointers. Strategy:

- Sweep all available memory, test each word for the *isPointer* property
- Unless the collector can (heuristically) decide that a word does not contain a pointer, assume that it is a pointer.
- If the collector knows the range of valid memory addresses, it can efficiently exclude, small integers, real numbers and strings

315

Reference Counting

- Used in many languages and applications:
Smalltalk, InterLisp, Java, Adobe PhotoShop
- Storage management runs concurrently intermixed with program execution.
- Invariant: the reference count of each node is equal to the number of pointers pointing to the node.
- Can do optimizations on pointers local to a routine to reduce the amount of reference counting.

317

Compacting and Non-compacting Collectors

- Once recyclable memory has been identified, there are two main options for managing this memory.
- Compacting collectors move all live objects into a block of contiguous memory. This leaves a large block of memory available for reuse.
- Non-compacting collectors leave the live objects in place, and (somehow) link all the recycled memory for later use.
- Compacting advantages
 - Little effort required to reclaim objects.
 - New allocations are easy with a large block of contiguous free memory.
 - Moving object together may improve cache performance.
- Compacting Disadvantages
 - All pointer access must be managed through the collector.
 - Moving large blocks of memory may impact memory performance.
 - May require twice as much memory for compacting.

316

Reference Counting - Advantages

- Memory management overhead is distributed throughout the computation. May yield smoother response time.
- Locality of reference no worse than application.
- Data from real languages suggests that few nodes are shared and many nodes have a short lifetime. Reference counting lets these nodes be reused as soon as they are free.
This may lead to better cache/virtual memory behavior.
This may allow *update in place* optimizations.
- Performance relatively independent of heap occupancy fraction.

318

Reference Counting - Disadvantages

- High processing cost. Every pointer assignment requires reference count updating.
- Correctness is difficult.
Missing even one reference counter update could lead to errors.
- Difficult to deal with embedded pointer assignments,
e.g. bulk assignment of a structure containing pointers.
- Reference count roughly doubles the size of every pointer.
- Reference counting schemes can not reclaim cyclical data structures, e.g. a doubly linked list.
- Need some strategy to deal with reference counter overflow.

319

Naive Mark Sweep Algorithm^a

- Basic algorithm:
 - Suspend program processing.
 - Start at all roots. The algorithm must be able to find **all** roots, e.g. even roots in registers.
 - Trace nodes reachable from the roots and mark them.
 - If trace reaches a node that is already marked, stop tracing on that path.
 - After marking is complete, make another sweep through the heap
unmarked nodes can be reclaimed.
reset mark on marked nodes for next sweep.
 - Restart program execution.
- Mark and Sweep is usually triggered asynchronously when the program makes a request for memory that cannot be satisfied.

^aUsed in original Lisp implementation.

320

Mark and Sweep Advantages


- Cycles handled easily.
- No overhead on pointer manipulation.
- No space overhead on pointers.
- Handles bulk assignment correctly.

Mark and Sweep - Disadvantages


- Stop - Start algorithm.
Unsuitable for interactive applications.
- Guaranteeing that all roots can be found is *hard*
- Need to disable collector in critical sections.
- Cost of mark/sweep passes is high especially in modern memory hierarchies.
- Time is proportional to size of heap, not number of active cells.
Algorithm will tend to *thrash* as heap occupancy increases.
- Without optimizations leads to highly fragmented memory.

321


Naive Mark and Sweep Example




322




323




324




325



326



327



328

Tri-color Marking

- A modern version of Mark/Sweep. This is an *on the fly* algorithm. Allocation and mutation can continue during collection.
- Start with three sets:
 - **White set** contains objects that can have their memory recycled.
 - **Black set** contains objects that have no pointers to objects in the white set but are not available for recycling.
 - **Gray set** contains all objects reachable from root references. Objects referenced from the grey set haven't been scanned yet. Grey set objects will eventually move to the Black set.
- Initially
 - The Black set is empty
 - The Grey set is initialized to all objects that can be reached *directly* from roots
 - The White set contains all objects not in the grey set.

329

- Algorithm
 - Objects only move from white to grey to black.
 - Pick an object (*victim*) from the Grey Set.
Move all objects in the White set that *victim* references *directly* to the Grey set.
Move object to the Black Set.
 - Iterate until the Grey set is empty.
Then all objects in the White set can be recycled.
- **Invariant** No object in the Black Set points directly to an Object in the White set.
Any allocation or mutation during collection must maintain this invariant.

330

Compacting Collectors

- Widely used due to superior performance over reference counting and mark/sweep.
- Uses two large regions of memory *fromSpace* and *toSpace*.
One region is active, containing data objects and new allocations
- Compacting collector algorithm
 - Suspend program execution.
 - Flip identity of two semi-spaces. $toSpace \leftrightarrow fromSpace$
 - Trace from roots all live data structures in *fromSpace*.
 - Copy live nodes to *toSpace* when the node is first visited.
 - Update all pointers on the fly to point to new location of node.
Assume *fromSpace* and *toSpace* pointers can be distinguished.
Requires mapping table of old/new pointer values.
 - *fromSpace* becomes inactive until the next collection
 - Restart program execution.

331

- Compacting Collection - Advantages
 - Good time performance. Efficient allocation.
 - Asymptotic complexity less than marks/sweep collectors.
 - Fragmentation is eliminated since live data is copied compactly into the *toSpace*.
 - Can handle bulk pointer assignment.
- Compacting Collection - Disadvantages
 - Must be able to identify and update *all* pointers to data objects being moved.
 - May degrade paging in virtual memory systems.
Collector may get more page faults than other techniques.
 - Requires potentially large mapping table during copying to update pointers.
 - Stop - start algorithm. Not suitable for interactive or real-time systems.
 - Requires double the address space of other techniques.

332